

Computer Vision

7/15/2025
Jing Jia





1

Computer Vision

A multidisciplinary field:


- Physics – light is an EM wave.
 - Optics – lenses, reflectance, etc.
 - Applied Mathematics (linear algebra, numerical methods, computational geometry)
 - Digital Signal Processing
 - Stochastics/Probability (dealing with uncertainty in estimation)
 - Computer Science (algorithm development)
-
- Computer Vision is a branch of AI
 - Computer Vision=Robot Vision
 - CV has overlap with Computer Graphics

What is Robot Vision, Computer Vision?

The goal of computer vision is to computationally represent visual imagery in order to enable a computer system to interpret and interact with the surrounding environment. Interpretation of imagery includes image segmentation, object/texture recognition, depth/distance estimation and motion estimation



Computer Vision in Real world



Applications of Computer Vision

- Recognition: Object, Face, Material
- Computer-assisted surgery (e.g. Frameless brain surgery)
- Autonomous navigation for Robotics
- Driverless Cars
- Content-based video indexing (e.g. find frames with a key person from thousands of hours of video)
- Hand Gesture Recognition (e.g. mouseless computers)
- Computer Graphics Rendering/ Movie Effects
- Computer Assisted Radiology/ Diagnosis (finding Tumors in X-rays)
- etc

CVPR 2013 – Expressive Visual Text-to-Speech



<https://www.youtube.com/watch?v=6QFIWP0ZfPY>

Ameca, humanoid robot from Engineered Arts



<https://www.youtube.com/watch?v=IPukuYb9xWw>

Amera and GPT3




<https://www.youtube.com/watch?v=fQ2JpkFBAOs>

<https://thispersondoesnotexist.com/>



Autonomous Driving



<https://www.youtube.com/watch?v=HS1wV9NMLr8>

Computer Vision

- What am I looking at? → Image Classification/Recognition
- What are all the things I see, name every image pixel? → Semantic Segmentation
- Where is the object in the 3D world? → 3D Reconstruction
- Where is it and where am I? → SLAM
- What should I do? → Reinforcement Learning
- How do I learn by watching you? → Imitation Learning




2

**Image
Formation**


Image Fundamentals

Definition of a Camera:

- A device that converts an optical signal into an electrical signal.



Pinhole Camera



<https://www.youtube.com/watch?v=jhBC39xZVnw>

Pinhole Camera in Nature: Dappled Sunlight




Fist Pinhole Camera



Perspective Projection

Notice that objects appear smaller as they move away from you



Principle of Image Acquisition

Pinhole Imaging Model:

- Simplest camera model, ignores field curvature and distortions.
- Real cameras use calibration to correct these distortions.
- Widely used due to simplicity and effectiveness.





Image Capture Process:

- Light passes through the camera lens onto a sensor chip.
- The sensor chip converts light information into a digital signal.
- The light-sensitive element on the sensor chip is crucial for this process.



Sensor chip

Principle of Image Acquisition



<https://www.youtube.com/watch?v=i85jZ76azZ8>

Pixels and Resolution

Pixels:

- Tiny square grids on a light sensor.
- The smaller the grid, the higher the pixel count.
- Higher pixel count results in higher image resolution.

Frame Rate (FPS):

- Frequency at which images are displayed, measured in frames per second.
- Human eye can barely detect lag above 20 fps.
- Higher frame rates are better for machine vision.



An image with a resolution of 31x22 pixels.

Color

Color Perception:


- Visual effect produced by the eye, brain, and individual experiences.
- Different frequencies of electromagnetic waves are seen as different colors.

RGB Color Model:

- Based on human eye sensitivity to yellow-green, green, and bluish violet light.
- Used in monitors for color display.

LAB Color Space:

- "L" represents luminance.
- "A" axis: red to green.
- "B" axis: yellow to blue.
- Designed to be close to human vision.




Color Theory



<https://www.youtube.com/watch?v=oSQj3meaJcQ>

LAB Color Space




The Lens

Lens Functionality:

- Refracts light to the sensor, determining image size and distance.

Focal Length:

- Distance from lens's main point to focal point.
- Determines image size, field of view, depth of field, and perspective.
- Longer focal length: larger image, smaller angle of view.
- Shorter focal length: smaller image, larger angle of view.



The Lens

Lens Aberration:


- Distances to the lens vary across the sensor, causing distortion.
- Fisheye effect (barrel distortion) at edges.
- Corrected using algorithms or distortion-free lenses.

Lens Comparison Example:

- Different focal lengths impact the visual range and image clarity.
- OpenMV example shows differences when 20cm away from the table.



Standard Lens



Wide-Angle Lens




Distortion-Free Lens



Telephoto Lens

Lens



<https://www.youtube.com/watch?v=EL9J3Km6wxI>



3

Matrix Representation

Vectors, Matrices

$$\text{Let } \vec{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}, \vec{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix}.$$

$$\vec{x}^T \vec{y} = \vec{x} \cdot \vec{y} = x_1 y_1 + x_2 y_2 + x_3 y_3 \text{ (A scalar!)} \quad (1)$$

$$\vec{x}^T = \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \quad (2)$$

$$\text{Matrix } A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}.$$

Matrix Multiplication

$$AB = C$$

$$c_{ij} = \sum_t a_{it} b_{tj}$$

c_{ij} is the inner product of the i th row of A and the j th column of B.

Note: Number of cols of A must equal Number of rows of B

Matrix multiplication properties:

- Associative $(AB)C = A(BC)$
- Distributive $A(B + C) = AB + AC$
- NOT Commutative $AB \neq BA$ in general

$$I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$
$$AI = A$$

Transpose

Transpose of A is denoted A^T . The i th row of A is the i th column of A^T .

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix}$$

A symmetric matrix is a matrix which equals its transpose

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 8 \end{bmatrix}$$

Linear Equations

Expressing a system of linear equations in matrix form

$$\begin{array}{l} 2x_1 + 3x_2 = 4 \\ 3x_1 + 4x_2 = 7 \end{array} \quad \rightarrow \quad Ax = b \quad \begin{array}{l} A = \begin{bmatrix} 2 & 3 \\ 3 & 4 \end{bmatrix} \\ x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \\ b = \begin{bmatrix} 4 \\ 7 \end{bmatrix} \end{array}$$

Linear Equations

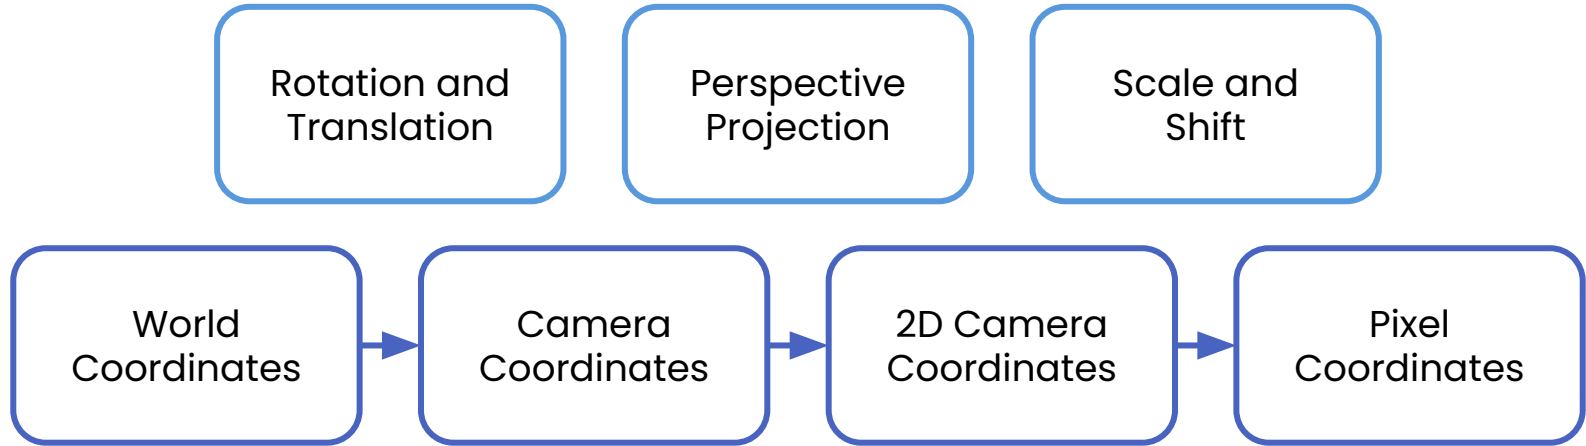
This idea of expressing equations with matrices is fundamentally important.

Ax can be viewed as a weighted sum of the columns of A .

$$\begin{bmatrix} 2 \\ 3 \end{bmatrix} x_1 + \begin{bmatrix} 3 \\ 4 \end{bmatrix} x_2 = \begin{bmatrix} 4 \\ 7 \end{bmatrix}$$

The system $Ax = b$ is solvable if and only if the vector b can be expressed as a linear combination of the columns of A .


Complete Image Formation Pipeline



The entire imaging pipeline, described by 3 steps, can be described by the product of three matrices... which is a single matrix

Perspective Projection: Similar Triangles

$$x = f \frac{X}{Z}$$
$$y = f \frac{Y}{Z}$$



Perspective Projection as Matrix Multiplication

It is desirable to have all of the mathematical relations in matrix form (easy concatenation of steps in the imaging pipeline, expression as a system of linear equations) The division by Z is a problem. We can't put the perspective projection in matrix form, these are nonlinear equations. We use **homogeneous coordinates or projective coordinates** to allow the perspective projection part of the imaging pipeline to be written as a matrix. When we use homogeneous coordinates ... **division is postponed**, Z remains in third component.

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

Homogeneous Coordinates

But these are homogeneous coordinates, so to get the final answer, divide by the third component.

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} \frac{\tilde{x}}{\tilde{w}} \\ \frac{\tilde{y}}{\tilde{w}} \end{bmatrix}$$

Shift and Scale (from image to pixel coordinates)

After perspective projection there is a scale and shift that transforms the projected values into pixels.

$$\begin{bmatrix} im_x \\ im_y \end{bmatrix} = \begin{bmatrix} 1/s_x & 0 & o_x \\ 0 & 1/s_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Why is a Shift Needed?

- Image pixels have a different x-y coordinate frame compared to camera coordinates.
- In an image, the origin (0,0) is typically at the upper left corner.
- In projected camera coordinates, the origin (0,0) is at the intersection of the principal axis with the image plane.
- This shift of the origin is accomplished using the O_x and O_y terms in the last column of the transformation matrix.

Why is a Scale Needed?

- The camera coordinate frame uses a unit of measure such as millimeters (mm).
- Pixels are unitless and require a conversion.
- Division by s_x and s_y (the horizontal and vertical size of the pixel in the measurement units, e.g., mm) is necessary to convert camera coordinates to pixel coordinates

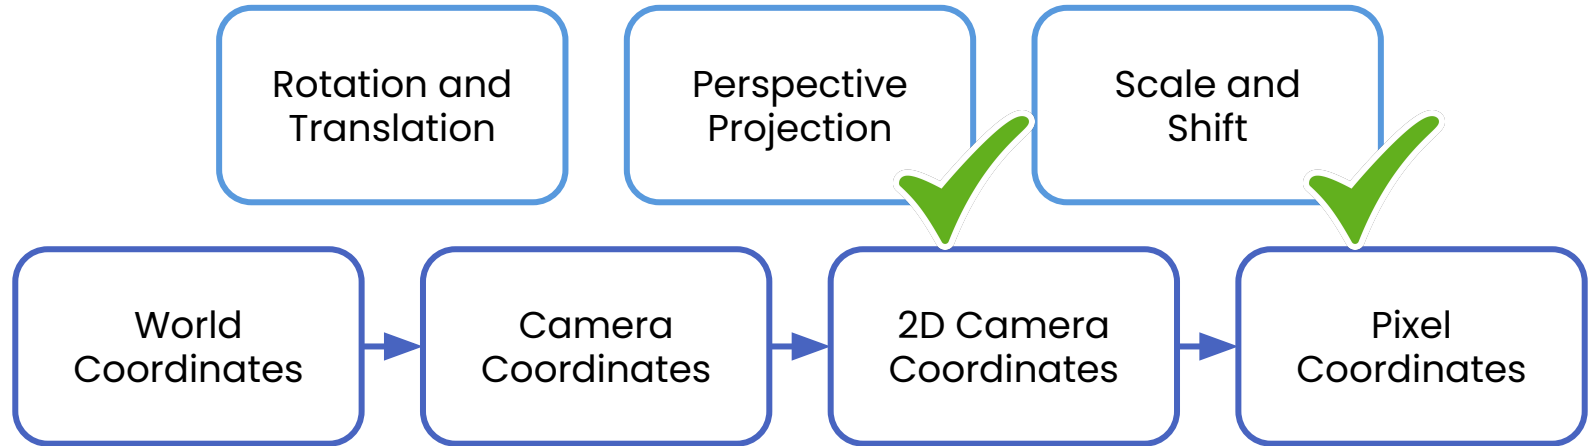
Perspective Project + Scale and Shift

Here we have denoted the pixel coordinates frame **im** by a superscript. We can put together the two matrices as follows:

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} = \begin{bmatrix} 1/s_x & 0 & o_x \\ 0 & 1/s_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} = \begin{bmatrix} f/s_x & 0 & o_x \\ 0 & f/s_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

Image Formation Pipeline now



Coordinate Frame Transformations

- Perspective Projection must be done in a camera coordinate frame
- Object Vertices are in a World Coordinate Frame (a reference frames).
- For example a cube has vertices
 - Bottom Face: $(0,0,0)$; $(1,0,0)$; $(1,1,0)$; $(0,1,0)$
 - Top Face: $(0,0,1)$; $(1,0,1)$; $(1,1,1)$; $(0,1,1)$
- These vertices are with respect to a fixed reference coordinate frame.

3D Coordinate Frame Transformations

3x4 transformation matrix, rotation followed by a translation

$$T = \left[\begin{array}{c|c} R & t \end{array} \right]$$

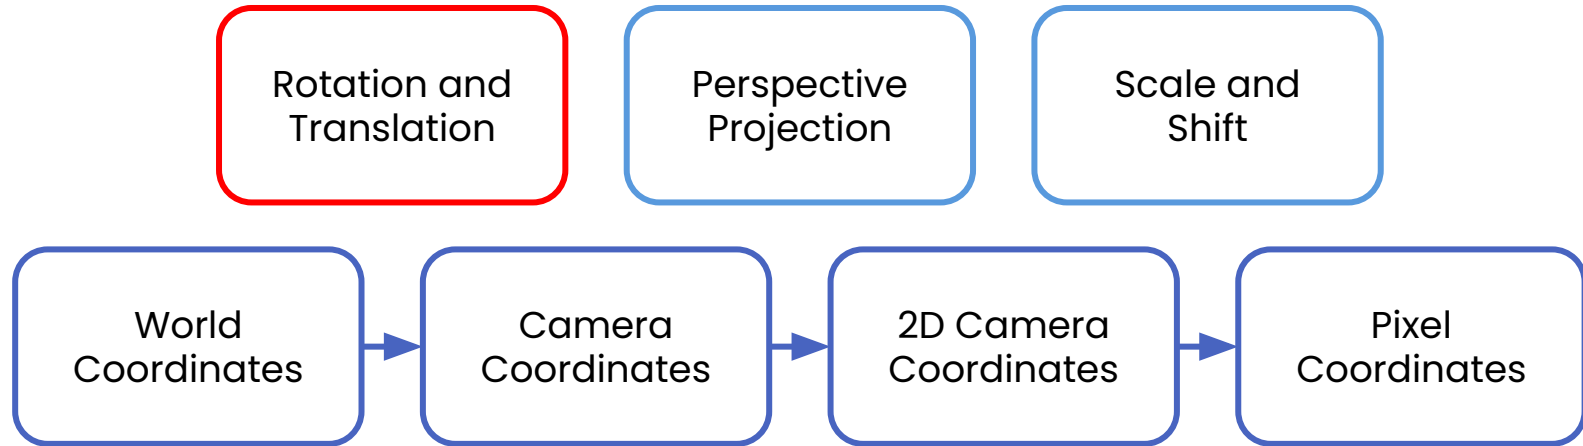


Image Formation Pipeline Complete

The complete image formation pipeline is given by,

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} = \begin{bmatrix} f/s_x & 0 & o_x \\ 0 & f/s_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} {}^C R & {}^C t \end{bmatrix} \begin{bmatrix} X_W \\ Y_W \\ Z_W \\ 1 \end{bmatrix}$$


Can be simplified to,

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} = M \begin{bmatrix} X_W \\ Y_W \\ Z_W \\ 1 \end{bmatrix}$$

Coordinate Frame Transformation

We will consider two ways to specify the coordinate frame transformation:

- 1) with a series of rotations followed by translations
- 2) using the geometric interpretation of the columns of the rotation matrix.



Transformation Process

Initial Alignment:

- Let A denote the reference frame and B denote the other frame.
- Coordinate frame B is initially aligned with A.

Rotation Steps:

- We rotate B about x-axis of A by angle α .
- Then rotate B about y-axis of A by angle β .
- Finally, rotate B about z-axis of A by angle θ .

Translation:

- The coordinate frame B is translated by 10 units along x-axis of A .

Coordinate Frame Transformation

Given a point P w.r.t the B coordinate frame we can convert it to that point w.r.t. the A coordinate frame as follows

$${}^A P = \begin{bmatrix} {}^A R_B & {}^A t_B \end{bmatrix} ({}^B P) = {}^A T_B ({}^B P)$$

$${}^A R_B = Rot(\hat{z}, \theta) Rot(\hat{y}, \beta) Rot(\hat{x}, \alpha)$$

Note the order of the concatenation of the rotation matrices.

The translation is given by

$${}^A t_B = \begin{bmatrix} 10 \\ 0 \\ 0 \end{bmatrix}$$



4

Object Recognition

What is Object Detection?

Classification




CAT

Classification
+ Locolation



CAT

Object Detection



CAT, DOG, DUCK

Instance
Segmentation



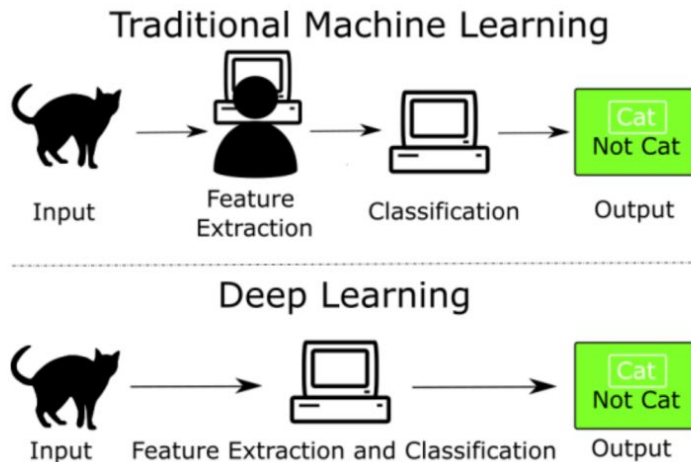
CAT, DOG, DUCK

Object Detection act as a combination of image classification and object localization. It takes an image as input and produces one or more bounding boxes with the class label attached to each bounding box.

Traditional Object Detection Era

Early Algorithms:

20 years ago, object detection relied on handcrafted features due to the lack of effective image representation.



Viola Jones Detectors

- **Development:** Created in 2001 by Paul Viola and Michael Jones.
- **Function:** Detects human faces in real-time using sliding windows to search for 'haar-like' features.
- **Speed Enhancements:**
 - Integral Image: Reduces computational complexity.
 - Adaboost Algorithm: Selects the most helpful features for face detection.
 - Detection Cascades: Multi-stage detection paradigm to reduce computational overhead.



(1)




(2)



(3)




(4)




Histogram of Oriented Gradients (HOG) Detector

- **Introduction:** Proposed in 2005 by N. Dalal and B. Triggs.
- **Function:** Uses blocks to constitute gradients from pixel intensity changes.
- **Application:** Widely used for pedestrian detection.
- **Size Adaptation:** Rescales the input image to detect objects of different sizes.



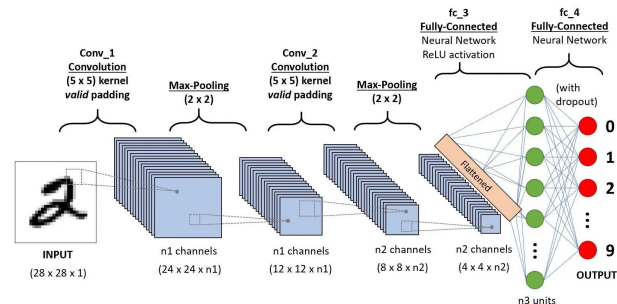
Deformable Part-based Model (DPM)

- **Proposed by:** P. Felzenszwalb in 2008, with improvements by R. Girshick.
- **Strategy:** Divides the detection of an object into its parts (e.g., car's window, body, wheels).
- **Components:** Root-filter and part-filters learned as latent variables.
- **Techniques for Accuracy:**
 - Multi-Instance Learning: Special case used for improving detection accuracy.
 - Additional Techniques: Hard negative mining, bounding box regression, context priming.
 - Cascade Architecture: Achieves acceleration without sacrificing accuracy.




Deep Learning Era

- **Performance Plateau:** Post-2010, performance of handcrafted features plateaued.
- **CNN:** In 2012, convolutional neural networks (CNNs) enabled robust and high-level feature representations.
- **Breakthrough:** 2014 saw the proposal of Regions with CNN features (RCNN) breaking the deadlocks of object detection.
- **Genres:** Object detection is grouped into two-stage detection and one-stage detection.



Why Object Detection?

- Robot navigation
- Medical diagnosis
- Security
- Industrial inspection and automation
- Human-computer interface
- Information retrieval



Applications of Object Recognition






5


K-Nearest Neighbors

Image Classification



An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```



Plane


We usually take image classification in computer vision as example

- Computer vision is the most successful field for DNNs
- Image classification is the core task of computer vision

Challenge: Semantic Gap

In computer, each pixel is represented with a number

- Computer views an image as a big grid of numbers




Challenge: Different Sizes, Shapes, Angles...



Machine Learning: Data -driven Approach

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images




A Naive Idea: Find Similar Images

```
def train(images, labels):  
    # Machine Learning!  
    return model
```



Memorize all data and labels

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```





Predict the label of the most similar training images

Nearest Neighbors in Training Image

Test Image



Training Image



How to Determine the Nearest?


P-norm is good metric:

p: Dimension (or pixel position)

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

10	20	24	17
8	80	109	12
12	16	178	170
4	32	233	112




46	12	14	1
82	43	19	121
12	10	0	30
2	32	22	108

Distance Metric


L1(Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

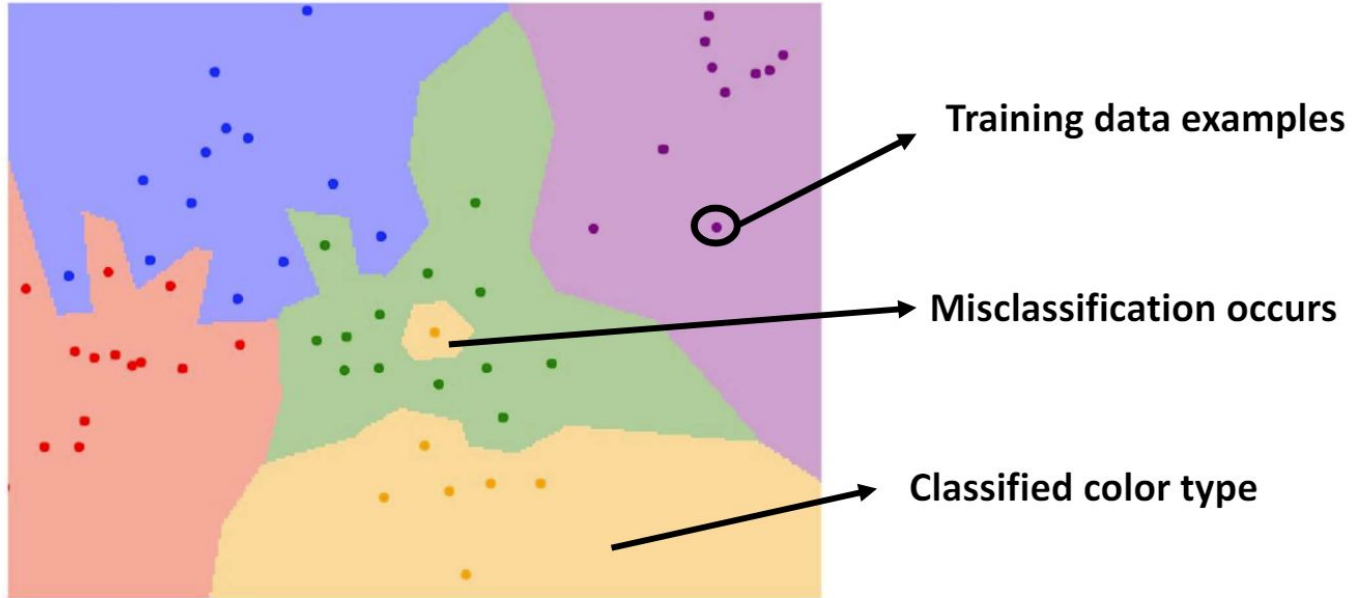


L2(Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



Disadvantage If Only Considering Nearest



Improvement: Using K Nearest Neighbors (KNN)

KNN: Take majority vote from K closest points



K = 1



K = 3



K = 5

Hyperparameter

When setting up KNN, you can choose two parameters:

- What is the best value of k to use for voting?
- What is the best distance to use for measuring?

These are hyperparameters, which are not adapted by the machine learning algorithm itself.

In DNN, many hyperparameters

- # of layers/neuron, learning rate, weight decay.....

Train, Validation and Test Dataset

- **Training set**


- a set of data used to discover predictive relationships
- simple words: help to learn the model

- **Test set**

- a set of data used to assess the strength and utility of a predictive relationship
- simple words: help to evaluate a learned model

- **Validation set**

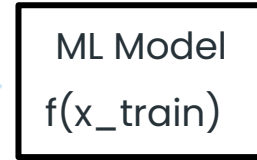
- a set of examples used to tune the hyperparameters



Summary of Non - parametric KNN

Training Data

0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9




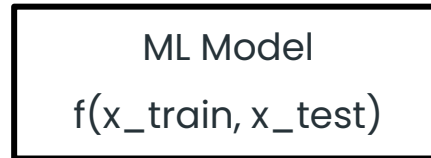
Pros:

Fast Training
Easy to implement

Cons:

Slow Testing
Store all training data

Test Data





6


Convolutional Neural Networks

History

PROC. OF THE IEEE, NOVEMBER 1998

Gradient-Based Learning Applied to Document Recognition

Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner



Yann LeCun

- Machine Learning
- Computer Vision
- Mobile Robotics
- Computational Neuroscience

Key Contributions:

- Optical Character Recognition & Computer Vision: Known for pioneering work using convolutional neural networks (CNNs).
- Often referred to as one of the "Godfathers of AI" and "Godfathers of Deep Learning."




Convolutional Network Demo from 1989



https://www.youtube.com/watch?v=FwFduRA_L6Q

ImageNet Challenge in 2012


ImageNet team managed to organize the ImageNet Large Scale Visual Recognition Challenge (ILSVRC). Competitors had to correctly classify images and detect different objects and scenes across a trimmed list of 1,000 ImageNet categories. Every year, the team that produced the model with the lowest error rate won




AlexNet

AlexNet didn't just win; it dominated.

AlexNet was unlike the other competitors. This new model demonstrated unparalleled performance on the largest image dataset of the time, ImageNet. This event made AlexNet the first widely acknowledged, successful application of deep learning. It caught people's attention with a 9.8 percentage point advantage over the nearest competitor



AlexNet



Deep Learning Image Classification...



**...with Convolutional
Neural Networks!!!**

<https://www.youtube.com/watch?v=HGwBXDKFk9I>

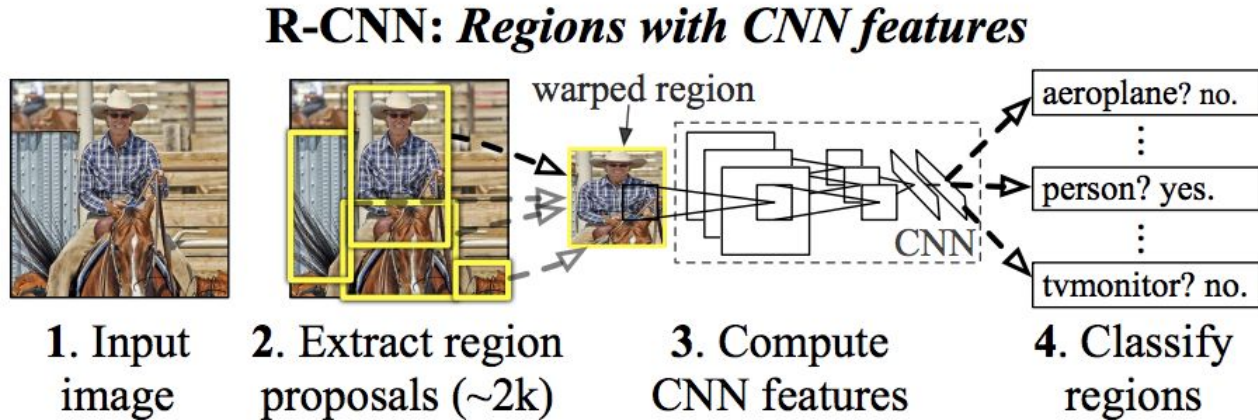


7

R-CNN

RCNN : Region Proposal + CNN

Region-based Convolutional Neural Network (R-CNN) is a type of deep learning architecture used for object detection in computer vision tasks. RCNN was one of the pioneering models that helped advance the object detection field by combining the power of convolutional neural networks and region-based approaches.



Selective Search

Selective Search is a **region proposal algorithm** for object detection tasks. It starts by over-segmenting the image based on intensity of the pixels. Selective Search then takes these oversegments as initial input and performs the following steps

1. Add all bounding boxes corresponding to segmented parts to the list of regional proposals
2. Group adjacent segments based on similarity
3. Go to step 1

At each iteration, larger segments are formed and added to the list of region proposals. Hence we create region proposals from smaller segments to larger segments in a bottom-up approach.


Selective Search

1. Generate initial sub-segmentation of input image using the method describe by Felzenszwalb et al in his paper "Efficient Graph-Based Image Segmentation".




Selective Search

2. Recursively combine the smaller similar regions into larger ones. We use Greedy algorithm to combine similar regions to make larger regions.



Selective Search

3. Use the segmented region proposals to generate candidate object locations



Feature Extraction

Once the region proposals are generated, approximately 2,000 regions are extracted and anisotropically warped to a consistent input size that the CNN expects (e.g., 224x224 pixels) and then it is passed through the CNN to extract features.


The output of the CNN is a high-dimensional feature vector representing the content of the region proposal.



Object Classification

The extracted feature vectors from the region proposals are fed into a separate machine learning classifier for each object class of interest. R-CNN typically uses Support Vector Machines (SVMs) for classification. For each class, a unique SVM is trained to determine whether or not the region proposal contains an instance of that class.

During training, positive samples are regions that contain an instance of the class. Negative samples are regions that do not.



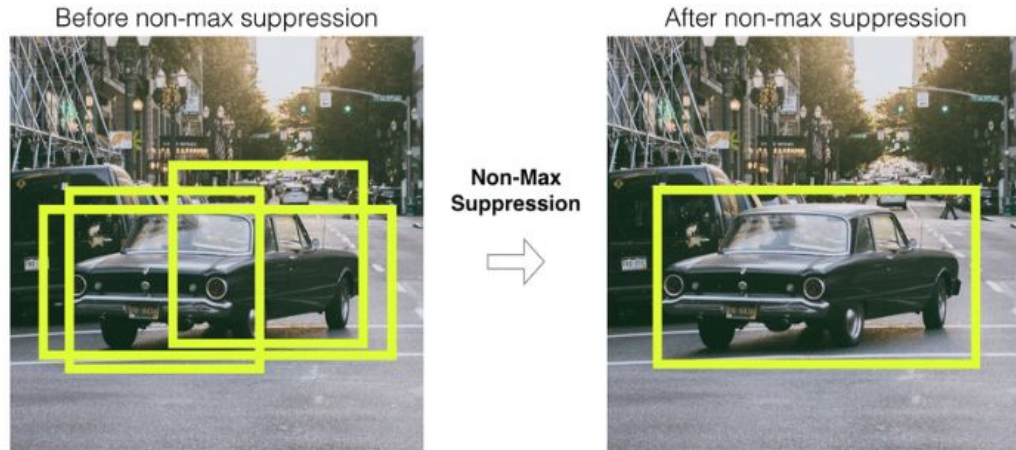
Support Vector Machines (SVMs)



https://www.youtube.com/watch?v=_YPScrckx28

Non-Maximum Suppression (NMS)

After classifying and getting bounding boxes for each region proposal, R-CNN applies non-maximum suppression to eliminate duplicate or highly overlapping bounding boxes. NMS ensures that only the most confident and non-overlapping bounding boxes are retained as final object detections.



R-CNN Pro and Cons

Pros

- Accurate Object Detection: Leverages region-based convolutional features for precise localization and recognition.
- Robustness to Object Variations: Handles different sizes, orientations, and scales, suitable for diverse objects and complex backgrounds.
- Flexibility: Adaptable to tasks, including instance segmentation and object tracking, by modifying the final network layers.

Cons

- Slow Inference: Sequential processing of region proposals results in latency, unsuitable for real-time applications.
- Overlapping Region Proposals: Generates multiple overlapping proposals, leading to redundant computation.
- Not End-to-End: Separate modules for region proposal and classification can result in suboptimal performance



8

Top Models


YOLO (You Only Look Once) Series

- **Architecture:** YOLO's architecture is inherently different from the region's proposal-based methods. Instead of proposing potential object locations and then classifying them, YOLO performs both tasks simultaneously. This concurrent approach is what gives YOLO its name and its speed.
- **Analysis:** Its standout feature is its speed. YOLO can process images in real-time, making it one of the fastest object detection models available. However, this design also means that YOLO might struggle with small objects or objects that are close together, as they might fall into the same grid cell.

YOLO Application

- **Traffic Management:** YOLO's real-time processing capability is invaluable in traffic management systems. For instance, in smart cities, YOLO is used to monitor intersections, detect traffic violations, and predict congestion. Its ability to swiftly detect vehicles, pedestrians, and even traffic signs aids in real-time decision-making, optimizing traffic flow and enhancing safety.
- **Real-time Surveillance:** Security systems leverage YOLO for intruder detection, unauthorized access, and monitoring restricted areas. The speed of YOLO ensures that security systems detect and address threats immediately.
- **Sports Analytics:** In sports, YOLO aids in player tracking, ball trajectory prediction, and game analysis. For example, in football, YOLO can track players' movements, analyze team formations, and even predict potential goal-scoring opportunities.

YOLO



<https://www.youtube.com/watch?v=QOC6vgnWnYo>

Detectron2

- **Architecture:** Detectron2 is the next-generation library developed by Facebook AI Research, succeeding Detectron and maskrcnn-benchmark. Built on the PyTorch framework, it offers a modular library rich with a variety of pre-trained models. Detectron2 provides state-of-the-art detection and segmentation algorithms, supporting numerous computer vision research projects and production applications.
- **Analysis:** Detectron2 is not just an incremental improvement over its predecessors, but a complete overhaul. It introduces new capabilities such as panoptic segmentation, Densepose, Cascade R-CNN, PointRend, DeepLab, ViTDet, MViTv2, and more. The library is both a research tool and a production-ready system. Its modular nature allows for easy customization, making it a preferred choice for researchers and developers alike.

Detectron2



<https://www.youtube.com/watch?v=GwQxaE8teRc>

EfficientDet

- **Architecture:** EfficientDet is an innovative object detection model that emphasizes efficiency without compromising on performance. It's built on the foundation of compound scaling – a method that uniformly scales the resolution, depth, and width of the network. This ensures that the model remains balanced and optimized across various device capabilities. One of the architectural innovations in EfficientDet is the introduction of a weighted bi-directional feature pyramid network (BiFPN). This BiFPN facilitates easy and rapid multiscale feature fusion, enhancing the model's ability to detect objects of varying sizes.
- **Analysis:** The BiFPN allows for faster feature fusion across different scales. Additionally, the compound scaling method ensures that all components of the model (backbone, feature network, and prediction networks) scale uniformly, maintaining a harmonious balance.

EfficientDet



<https://www.youtube.com/watch?v=AEgkLRZnm3I>



9

AprilTags

AprilTags


AprilTag is a visual fiducial system, useful for a wide variety of tasks including augmented reality, robotics, and camera calibration. Targets can be created from an ordinary printer, and the AprilTag detection software computes the precise 3D position, orientation, and identity of the tags relative to the camera. The AprilTag library is implemented in C with no external dependencies. It is designed to be easily included in other applications, as well as be portable to embedded devices. Real-time performance can be achieved even on cell-phone grade processors.




<https://april.eecs.umich.edu/software/apriltag>

AprilTags


AprilTags are conceptually similar to QR Codes, in that they are a type of two-dimensional bar code. However, they are designed to encode far smaller data payloads (between 4 and 12 bits), allowing them to be detected more robustly and from longer ranges. Further, they are designed for high localization accuracy—you can compute the precise 3D position of the AprilTag with respect to the camera.




Tag36h11




TagStandard41h12




TagStandard52h13



TagCircle21h7



TagCircle49h12




TagCustom48h12

VisionPortal from FIRST Tech Challenge


VisionPortal is a comprehensive new interface for vision processing. VisionPortal offers key capabilities of AprilTag and EasyOpenCV, along with TensorFlow Object Detection (TFOD)

- **AprilTag**
 - Detects an ID code.
 - Estimates pose (position and rotation).
- **TensorFlow Object Detection (TFOD)**
 - Recognizes 3D physical objects
 - Defines a "bounding box" and estimates heading.



HuskyLens

HuskyLens AI Vision Sensor can be used in the FIRST Tech Challenge Game. It offers Tag Recognition function which can detect tags, and learn, recognize, track specified tags.



<https://www.youtube.com/watch?v=Kv1HiN7nmg4>





10

OpenCV

What Is OpenCV?

- OpenCV: An open-source computer vision library available at [SourceForge.net](https://sourceforge.net/projects/opencvlibrary/).
- Languages: Written in C and C++, with active development for interfaces in Python, Ruby, Matlab, and others.
- Platforms: Runs on Linux, Windows, and Mac OS X.

Design and Efficiency

- Computational Efficiency: Designed for real-time applications.
- Optimized C Code: Takes advantage of multicore processors.

OpenCV Goals

- Vision Infrastructure: Provides a simple-to-use framework for building sophisticated vision applications quickly.
- Functionality: Contains over 500 functions for various vision tasks:
 - Factory product inspection
 - Medical imaging
 - Security
 - User interface
 - Camera calibration
 - Stereo vision
 - Robotics

Core Operations

Function	Purpose
<code>imread(image_path, flag)</code>	This method is used to read an image from its path
<code>imshow(window_name, image)</code>	It is used to show the image in the window.
<code>imwrite(filename, image)</code>	This method is used to write or save an image using OpenCV.

Reading Images

Image Reading Modes	Description
<code>cv2.IMREAD_COLOR</code>	Read colored images
<code>cv2.IMREAD_GRAYSCALE</code>	Read Grayscale images
<code>cv2.IMREAD_UNCHANGED</code>	Read images with alpha channel

Drawing Shapes and Text on Images

Syntax	Purpose
<code>line(image, start_coordinates, end_coordinates, color_in_bgr, line_thickness)</code>	By using this function we can create a line in the image from start coordinates to end coordinates with a certain thickness which can be mentioned specifically as well.
<code>rectangle(image, top_left_vertex_coordinates, lower_right_vertex_coordinates, color_in_bgr, thickness)</code>	This function is used to create a box with a certain thickness and color which can be specified as well.
<code>circle(image, center_coordinates, radius, color, thickness)</code>	It is used to draw a circle whose centre and radius length is given with a certain thickness and the colour of the strokes of the circle.

Drawing Shapes and Text on Images

Syntax	Purpose
<code>polylines(image, [pts], isClosed, color, thickness)</code>	It is used to draw a polygon on any image whose vertex coordinates are provided.
<code>putText(image, 'TextContent', 'text_starting_point_coordinates', 'fontToBeUsed', 'font_size', 'text_color', 'text_thickness', 'line_type')</code>	It is used to write some text on the image loaded.

Arithmetic Operations on Images

Function	Uses
<code>add(image1, image2)</code>	This function is used to add two images.
<code>subtract(image1, image2)</code>	This function is used to subtract two images.
<code>addWeighted(image1, weight1, image2, weight2, gammaValue)</code>	This is also known as Alpha Blending. This is nothing but a weighted blending process of two images.

Arithmetic Operations on Images

Function	Uses
<code>bitwise_and(image1, image2, destination, mask)</code>	This performs bitwise and logical operations between two images.
<code>bitwise_or(image1, image2, destination, mask)</code>	This performs bitwise or logical operations between two images.
<code>bitwise_not(image, destination, mask)</code>	This performs bitwise not logical operations between an image and a mask.
<code>bitwise_xor(image1, image2, destination, mask)</code>	This performs bitwise xor logical operations between two images.

Geometric Transformations on Image

Function	Purpose
<code>resize(image, width, height, interpolation_method)</code>	<p>By using this method we can resize the image at hand for specific width and height.</p> <p>Some of the methods that can be used for interpolation are as below:</p> <ul style="list-style-type: none">• <code>cv2.INTER_AREA</code>• <code>cv2.INTER_CUBIC</code>• <code>cv2.INTER_LINEAR</code>

Geometric Transformations on Image

Function	Purpose
<pre>M = cv2.getRotationMatrix2D(center, angle, scale) warpAffine(image, M, (width, height))</pre>	<p>Image rotation is a common image processing routine with applications in matching, alignment, and other image-based algorithms, in image rotation the image is rotated by a definite angle. We can also use <code>cv2.rotate()</code> function for image rotation.</p>

Geometric Transformations on Image


Function	Purpose
<pre>matrix = cv2.getPerspectiveTransform(src, dst) warpPerspective(image, matrix, dsize)</pre>	<p>In Perspective Transformation, we can change the perspective of a given image or video for getting better insights into the required information. We need to provide the points on the image from which want to gather information by changing the perspective. We also need to provide the points inside which we want to display our image.</p>

Edge Detection

Edge Detection

Basic Idea

- * Look for a neighborhood with strong signs of change



Edge/Line Detection (Features)

Operation	Uses
<code>cv2.Canny(image, T_lower, T_upper, aperture_size, L2Gradient)</code>	The <u>Canny Edge Detection</u> is an algorithm used for edge detection. It reduces noise using Gaussian Smoothing and computes image gradient using The Sobel filter.
<code>cv2.HoughLines(edges, 1, np.pi/180, 200)</code>	The <u>Hough Transform</u> is a method that is used in image processing to detect any shape if that shape can be represented in mathematical form. It can detect the shape even if it is broken or distorted a little bit.

Edge/Line Detection (Features)

Operation	Uses
<code>cv2.HoughCircles(image, cv2.HOUGH_GRADIENT, 1, 20, param1 = 50, param2 = 30, minRadius = 1, maxRadius = 40)</code>	<u>Circle detection</u> finds a variety of uses in biomedical applications, ranging from iris detection to white blood cell segmentation.
<code>cv2.cornerHarris(src, dest, blockSize, kSize, freeParameter, borderType)</code>	<u>Harris Corner detection</u> algorithm was developed to identify the internal corners of an image. The corners of an image are basically identified as the regions in which there are variations in the large intensity of the gradient in all possible dimensions and directions.

Changing the Colorspace of Images

```
cv2.cvtColor(image, conversion_scale)
```

Some of the commonly used ways in which color coding of the images are changed is as shown below:

- cv2.COLOR_BGR2GRAY
- cv2.COLOR_BGR2HSV
- cv2.COLOR_BGR2LAB
- cv2.COLOR_BGR2YCrCb

```
hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
lower_blue = np.array([110,50,50])
upper_blue = np.array([130,255,255])
cv2.inRange(hsv, lower_blue, upper_blue)
```

OpenCV provides functions for the detection of a specific color in live-stream video content. A video is composed of infinite frames at different time instants.

Working With Videos

```
cv2.VideoCapture("file_name.mp4")
```

To capture a video using OpenCV, we need to create a VideoCapture object. VideoCapture has the device index or the name of a video file. The device index is just the number to specify which camera. If we pass 0 then it is for the first camera, 1 for the second camera.

```
cv2.VideoCapture(File_path)
```

```
cv2.read()
```

```
cv2.imwrite(filename, img[, params])
```

Using OpenCV, techniques such as image scanning, and face recognition can be accomplished quite easily. We can

extract images from videos as well using the OpenCV module along with the os module.

11


Raspberry Pi Camera

Raspberry Pi Camera Module



Raspberry Pi Camera Module

- Connect the Camera
- Please power off your Pi before plugin the Camera
- Please it's at the correct orientation



Test Camera

To test camera, use following command in terminal:

```
rpicom-hello
```

Python Packages for PiCamera

Create a new virtual environment with default packages including picamera2:

```
python -m venv --system-site-packages ~/env/opencv
```

Install opencv:

```
source ~/env/opencv/bin/activate
```

```
pip install opencv-python
```

Capture & Display Image

```
from picamera2 import Picamera2
```

```
import cv2
```

```
picam2 = Picamera2()
```

```
picam2.start()
```

```
image = picam2.capture_array() # Capture an image with RGB encoding
```

```
image_bgr = cv2.cvtColor(image, cv2.COLOR_RGB2BGR) # Convert from RGB to BGR
```

```
cv2.imshow("Captured Image", image_bgr)
```

```
cv2.waitKey(0)
```